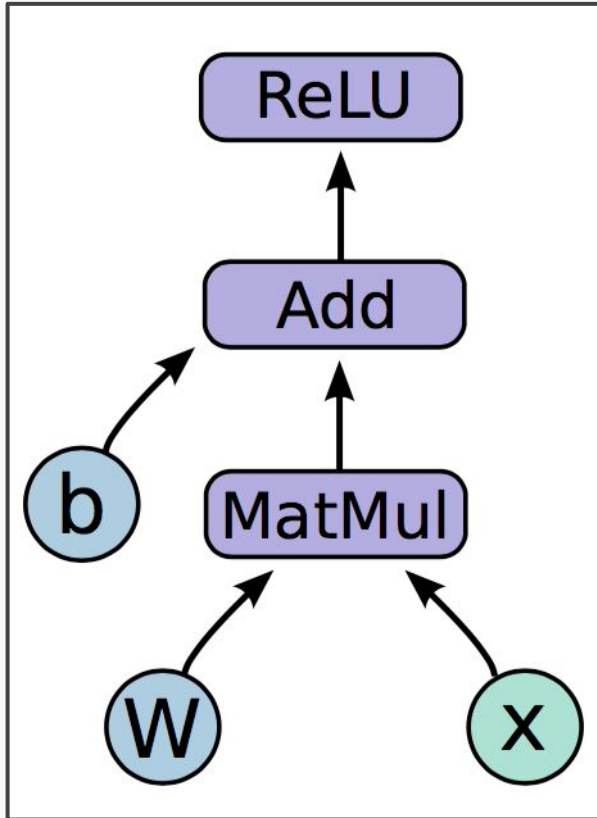


# Programming model

Big idea: express a numeric computation as a **graph**.

- Graph nodes are **operations** which have any number of inputs and outputs
- Graph edges are **tensors** which flow between nodes


$$h = \text{ReLU}(Wx + b)$$



$$h = \text{ReLU}(Wx + b)$$

**Variables** are stateful nodes which output their current value.  
State is retained across multiple executions of a graph

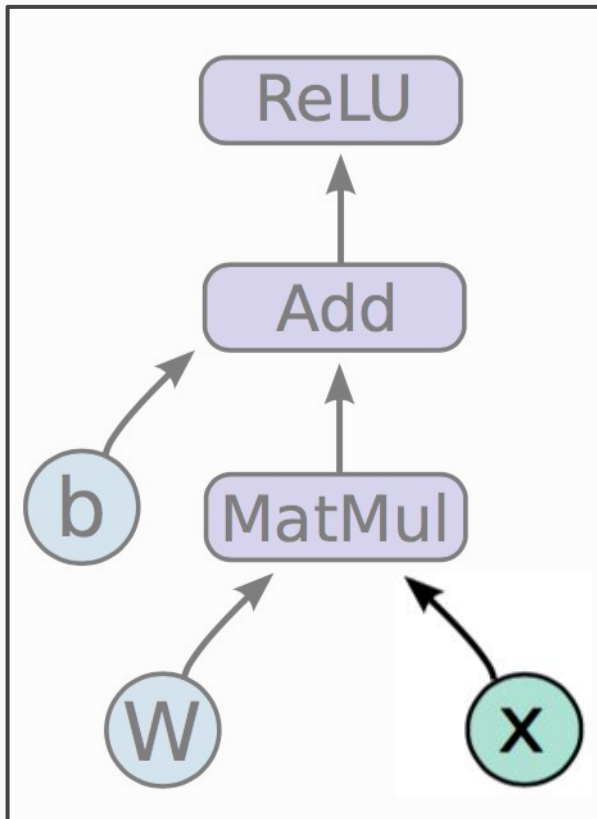
(mostly parameters)



$$h = \text{ReLU}(Wx + b)$$

**Placeholders** are nodes whose value is fed in at execution time

(inputs, labels, ...)



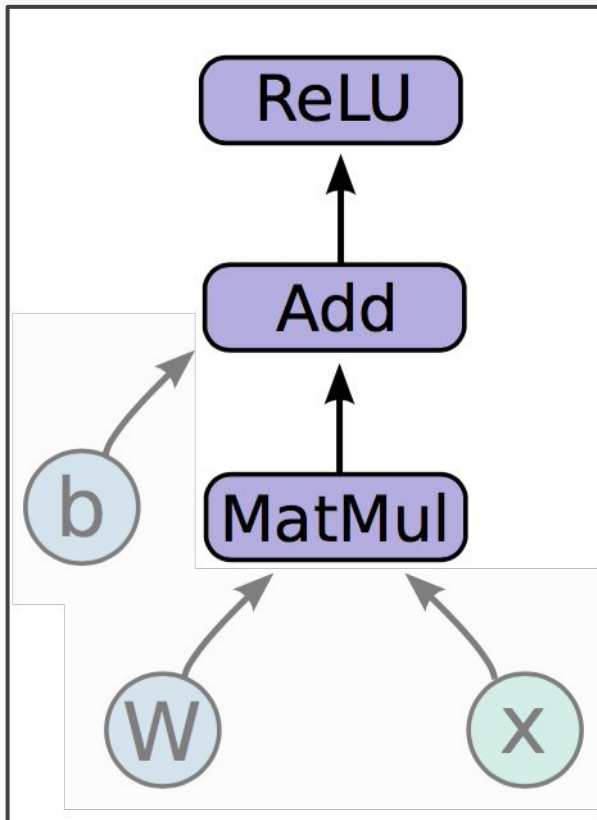
$$h = \text{ReLU}(Wx + b)$$

## Mathematical operations:

**MatMul:** Multiply two matrix values.

**Add:** Add elementwise (with broadcasting).

**ReLU:** Activate with elementwise rectified linear function.



## In code,

1. Create weights, including initialization

$$W \sim \text{Uniform}(-1, 1); b = 0$$

2. Create input placeholder  $x$   
 $m * 784$  input matrix

3. Build flow graph

```
import tensorflow as tf
```


```
b = tf.Variable(tf.zeros((100,)))
```

```
W = tf.Variable(tf.random_uniform((784, 100), -1, 1))
```

```
x = tf.placeholder(tf.float32, (100, 784))
```

```
h = tf.nn.relu(tf.matmul(x, W) + b)
```

$$h = \text{ReLU}(Wx + b)$$




# But where is the graph?

New nodes are automatically built into the underlying graph!  
`tf.get_default_graph().get_operations():`

zeros/shape  
zeros/Const  
zeros  
Variable  
Variable/Assign  
Variable/read  
random\_uniform/shape  
random\_uniform/min  
random\_uniform/max  
random\_uniform/RandomUniform

random\_uniform/sub  
random\_uniform/mul  
random\_uniform  
Variable\_1  
Variable\_1/Assign  
Variable\_1/read  
Placeholder  
MatMul  
add  
**Relu == h**


h refers to an op!



# How do we run it?

So far we have defined a **graph**.

We can deploy this graph with a **session**:  
a binding to a particular execution  
context (e.g. CPU, GPU)





# Getting output

```
sess.run(fetches, feeds)
```

**Fetches:** List of graph nodes.

Return the outputs of these nodes.

**Feeds:** Dictionary mapping from graph nodes to concrete values. Specifies the value of each graph node given in the dictionary.

```
import numpy as np
import tensorflow as tf
```

```
b = tf.Variable(tf.zeros((100,)))
W = tf.Variable(tf.random_uniform((784, 100),
                                  -1, 1))
```

```
x = tf.placeholder(tf.float32, (100, 784))
h = tf.nn.relu(tf.matmul(x, W) + b)
```

```
sess = tf.Session()
sess.run(tf.initialize_all_variables())
sess.run(h, {x: np.random.random(100, 784)}))
```

# So what have we covered so far?

We first built a **graph** using **variables** and **placeholders**

We then deployed the graph onto a **session**, which is the **execution environment**

Next we will see how to **train** the **model**

# How do we define the loss?

Use **placeholder** for **labels**

Build loss node using labels and **prediction**

```
prediction = tf.nn.softmax(...) #Output of neural network
label = tf.placeholder(tf.float32, [100, 10])

cross_entropy = -tf.reduce_sum(label * tf.log(prediction), axis=1)
```

# How do we compute Gradients?

```
train_step = tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)
```

- `tf.train.GradientDescentOptimizer` is an **Optimizer** object
- `tf.train.GradientDescentOptimizer(lr).minimize(cross_entropy)` adds optimization **operation** to computation graph

TensorFlow graph **nodes** have **attached gradient operations**

Gradient with respect to **parameters** computed with **backpropagation**

*...automatically*

# Creating the train\_step op


```
prediction = tf.nn.softmax(...)
label = tf.placeholder(tf.float32, [None, 10])

cross_entropy = tf.reduce_mean(-tf.reduce_sum(label * tf.log(prediction),
reduction_indices=[1]))

train_step = tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)
```



# Variable sharing



# Variable sharing: naive way

```
variables_dict = {  
    "weights": tf.Variable(tf.random_normal([784, 100]),  
                           name="weights"),  
    "biases": tf.Variable(tf.zeros([100]), name="biases")  
}
```

Not good for encapsulation!



# What's in a Name?

`tf.variable_scope()` provides simple name-spacing to avoid clashes

`tf.get_variable()` creates/accesses variables from within a variable scope

```
with tf.variable_scope("foo"):
    v = tf.get_variable("v", shape=[1]) # v.name == "foo/v:0"

with tf.variable_scope("foo", reuse=True):
    v1 = tf.get_variable("v") # Shared variable found!

with tf.variable_scope("foo", reuse=False):
    v1 = tf.get_variable("v") # CRASH foo/v:0 already exists!
```

# In Summary:

1. Build a graph
  - a. Feedforward / Prediction
  - b. Optimization (gradients and train\_step operation)
2. Initialize a session
3. Train with `session.run(train_step, feed_dict)`